

# Using the NIH HPC Storage Systems Effectively

Tim Miller

[btmiller@hpc.nih.gov](mailto:btmiller@hpc.nih.gov)

<https://hpc.nih.gov>

# Motivation

- Often get e-mail complaining about slow access to directories
- File storage systems are often over-saturated, leading to system problems that take staff time to resolve
- Users' work is often negatively impacted (sometimes without their knowledge).

# Presentation overview/outline

- Overview of HPC systems storage
  - Different areas (/home, /data, /scratch, object store)
  - Quotas and quota increases
  - Snapshots
- Understanding input and output (I/O)
  - I/O patterns
  - Introduction to data and metadata
  - A brief look at HPC storage systems
- Putting it all together – using storage effectively
  - Profiling and benchmarking your application's I/O usage
  - Understanding when you're generating too much I/O on the system

# Presentation overview/outline

- **Overview of HPC systems storage**
  - Different areas (/home, /data, /scratch, object store)
  - Quotas and quota increases
  - Snapshots
- Understanding input and output (I/O)
  - I/O patterns
  - Introduction to data and metadata
  - A brief look at HPC storage systems
- Putting it all together – using storage effectively
  - Profiling and benchmarking your application's I/O usage
  - Understanding when you're generating too much I/O on the system

# The NIH HPC filesystems

## Summary of file storage options

	Location	Creation	Backups	Space	Available from
/home	network (NFS)	with Helix account	yes	16 GB default quota	B,C,H
/lscratch (nodes)	local	created by user job	no	~850 GB shared	C
/scratch	network (NFS)	created by user	no	100 TB shared	B,H,C
/data	network (GPFS/NFS)	with Biowulf account	no	100 GB default quota	B,C,H

H = helix, B = biowulf login node, C = biowulf compute nodes

- /home is small – only filesystem backed up to tape – **No shared areas or quota increases!**
  - Snapshots and off-site tape back-ups
- /data - in practice you will keep most of your working files here
  - Quota increases available with justified need
  - Shared data directories available on request
  - Snapshots available (less frequent than /home)
  - ACLs
- /scratch – shared area for **temporary** data
- /lscratch – compute node local scratch; allocated with batch jobs

# General best practices

BAD	GOOD
Submitting a swarm without knowing how much data it will generate	Run a single job, sum up the output and tmp files, and figure out if you have enough space before submitting the swarm
Directory with 1 million files	Directories with < 5,000 files
100 jobs all reading the same 50 GB file over and over from /data/\$USER/	Use /lscratch instead, copy the file there, and have each job access the file on local disk
100 jobs all writing and deleting large numbers of small temporary files	Use <del>/scratch</del> instead, have all tmp files written to local disk <b>/lscratch</b>
Each collaborator having a copy of data on Biowulf	Ask for a shared area and keep shared files there to minimize duplication
Use Biowulf storage for archiving	Move unused or old data back to you local system

We'll be talking a lot about the “whys” behind some of these rules!

# The HPC Object Store

- “Web Scale” storage
  - Highly reliable (dispersed over multiple sites)
  - Easy to expand (just add more disks)
  - Accessed via simple list, put, get, delete semantics (examples forthcoming)
- Different from file based storage systems
  - Objects are accessed by **NAME**, not **PATH**
  - Completely flat name space
  - No concept of directories, but “/” is a valid character in object names
  - Data and metadata are stored together with the object (sometimes true in file storage systems as well)
- More info: <https://hpc.nih.gov/storage/object.html>
- Unless otherwise noted, the rest of this class deals with **file** (not object) storage. Separate object storage class!

# Use cases for object storage

- Read-intensive workloads
  - Object storage is much more efficient at reading than writing.
  - An **entire** object has to be re-written for each change
    - Computationally expensive to process and disperse the data
    - Lots of over-writing
- Static data
  - Related to the above
  - Data that doesn't change often, but still used
  - E.g. reference genomics files





# Understanding your disk quota

```
[teacher@helix ~]$ checkquota
```

Mount	Used	Quota	Percent	Files	Limit	Percent
/data:	6.8 GB	500.0 GB	1.37%	1695	31457280	0.01%
/home:	5.2 GB	16.0 GB	32.69%	1123	n/a	0.00%

ObjectStore Vaults	Used	Quota	Percent	Files	Limit	Percent
teacher:	5.7 MB	465.7 GB	0.00%	n/a	n/a	0.00%

- Use the checkquota command
  - Shows all directories you have access to.
- You can also get this information from your user dashboard
- Some storage systems have limits on the number of files – please keep these in mind (more on this later).

# Quota increases: on the user dashboard!

## User Dashboard

*last page refresh: 2020-05-19 17:11:45 EDT*

*page expires: 2020-05-19 19:11:12 EDT*

Accounts

Disk Usage

Job Info

Usage Report

### Diskspace Usage

*last updated: 2020-02-06 16:00:05 EDT*

/data/ (gs10)



27.9 GB / 500.0 GB owner:

/data/teacher (gs4)



6.8 GB / 500.0 GB [request quota increase](#)

/data/ (gs5)



8.6 GB / 100.0 GB owner:

/home/teacher



5.2 GB / 16.0 GB

/scratch/teacher



0.0 KB / 10.0 TB

# Quota increases: on the user dashboard!

## User Dashboard

*last page refresh: 2020-05-19 17:11:45 EDT*

*page expires: 2020-05-19 19:11:12 EDT*

Accounts

Disk Usage

Job Info

Usage Report

### Diskspace Usage

*last updated: 2020-02-06 16:00:05 EDT*

/data/ (gs10)



27.9 GB / 500.0 GB owner:

/data/teacher (gs4)



6.8 GB / 500.0 GB

[request quota increase](#)

/data/ (gs5)



8.6 GB / 100.0 GB owner:

/home/teacher



5.2 GB / 16.0 GB

/scratch/teacher



0.0 KB / 10.0 TB

Can only request quota increases on  
data directories you own!

# Quota increases: on the user dashboard!

- Only two things to fill in now!
  - Amount of space
  - **Justification**
- Also need to accept storage policies
  - NO PHI/PII
  - NO BACKUPS!

## HPC Disk Storage Request

This form should be completed by NIH HPC users who require additional storage space in their /data area.

There are no time limits or other restrictions placed on the use of data directory space on the Helix and Biowulf systems, but please use the space responsibly; even hundreds of terabytes won't last forever if files are never deleted. **Disk space on Helix and Biowulf should never be used as archival storage.**

### Location and User Information

Location: /data/ (gs2)  
Current Quota: 500 GB  
Name:   
Helx/Biowulf username:   
IC: NHLBI  
Telephone:   
Email: @nhlbi.nih.gov  
Principal Investigator:

### Space and Justification

Additional space required:  example: 250 GB

How this space requirement was estimated. Example: The input data for a single run of xx program is 100 MB. The output data is about 200 MB, and each run requires 100 MB of temp space. I expect to have about 600 such runs, so will need 400MB\*600 = 240 GB.

### Policies

Data directories are not backed up to tape. Irreplaceable data should be backed up to your local storage ([More info](#)). Personally Identifiable Information (PII) or Protected Health Information (PHI) data cannot be stored anywhere on the NIH HPC systems ([More info](#)).

Please confirm that you are aware of these policies: ☐ Agree

# Using space efficiently

- Instead of running to the quota request form, think about whether you can be using space more efficiently.
  - Move files back to your local computing infrastructure when you're done processing
  - Delete any raw data or intermediate files that you're sure you won't need again (or that are backed up elsewhere)
  - Compress (gzip, bzip2, etc.) files when not in active use (note: not all files compress well).
    - FASTQ files should always be compressed
    - Molecular dynamics binary trajectories generally don't compress much
- Storage space on the NIH HPC file systems is NOT to be used for archiving. However, we provide a time-limited archive for large data sets using the object store.

# Shared data directories

- Requested from [https://hpc.nih.gov/dashboard/shared\\_data\\_request.html](https://hpc.nih.gov/dashboard/shared_data_request.html)
  - Redirects to the user dashboard.
- A new group will be created (group name must begin with a capital letter).
  - Group members must be specified
- Justification for the storage request must be given as with a personal data directory quota increase.
- Don't open your personal data directory to world access!
  - Shared group directories should only be accessible by the group that owns them.
- Requestor becomes “group owner”
  - The only one who can request a quota increase
  - The only one who can request users be added to or removed from the group

# Shared data directories and permissions

```
[teacher@helix ~]$ ls -ld /data/SomeLab/  
drwxrws---+ 15 user SomeLab 32768 Apr 14 16:14 /data/SomeLab/
```

Group owner – the **ONLY** user  
allowed to request quota  
increases for the group.

- New directories created under the top level may **NOT** have the SGID bit set.
- Users can override group ownership and permissions of directories that they create within the shared area.
  - Some applications do this without informing the user.
- Coordination and care among group members is needed.
- Users in groups are responsible for maintaining correct permissions!

# Shared data directories and permissions

```
[teacher@helix ~]$ ls -ld /data/SomeLab/  
drwxrws---+ 15 user SomeLab 32768 Apr 14 16:14 /data/SomeLab/
```

Group name – same as the  
shared data directory name.

- New directories created under the top level may **NOT** have the SGID bit set.
- Users can override group ownership and permissions of directories that they create within the shared area.
  - Some applications do this without informing the user.
- Coordination and care among group members is needed.
- Users in groups are responsible for maintaining correct permissions!



# Shared data directories and permissions

```
[teacher@helix ~]$ ls -ld /data/SomeLab/  
drwxrws---+ 15 user SomeLab 32768 Apr 14 16:14 /data/SomeLab/
```

Permissions – all group members can write at the top level. The SGID bit (“s” in the group execute permissions) indicates all files created in this directory will have a group ownership of SomeLab.

- New directories created under the top level may **NOT** have the SGID bit set.
- Users can override group ownership and permissions of directories that they create within the shared area.
  - Some applications do this without informing the user.
- Coordination and care among group members is needed.
  - Set umask to 007 so that all files/directories created become group writeable
  - Some users request that their primary group be changed to the shared group
- Bonus: what does the “+” sign after the permissions mean?

# A few notes about /scratch

- /scratch is a **network accessed, global** /scratch directory
  - The same scratch directory is available on helix, biowulf, and all of the compute nodes.
  - Files deleted if not accessed for 7 days or when /scratch gets full!
- Good use-cases for /scratch
  - A temporary means of sharing data with a colleague who has a HPC account
  - Storing lightly accessed temporary files that must be accessed from multiple nodes.
- Bad use-cases for /scratch
  - Storing application temporary files that are only accessed from a single host (use /lscratch instead).
  - Storing anything that must be read and written frequently (use /data or /lscratch instead).
  - Storing valuable, hard to reproduce data (no back-ups or snapshots; use /home or /data as appropriate).

# Using /lscratch

- /lscratch = local scratch space on a compute node
- Allocated as a generic resource:
  - sbatch ... --gres=lscratch:100 ← allocates 100 GB
  - swarm -f swarmfile --gres=lscratch:100 ...
  - Ibid for sinteractive
  - At /lscratch/\$SLURM\_JOBID
- Benefits
  - Local to node, no network traffic
  - Fewer users sharing it
  - Nodes have fast solid-state disks



# When to use /lscratch (and when not)

- Use /lscratch when...
  - Many jobs will be independently reading in the same data file.
  - Many jobs will be independently writing out output files.
  - Many jobs will be writing out a lot of independent temporary files
  - Jobs will be doing large amounts of random I/O (more on this later)
  - **Swarms that read/write a lot of I/O should almost always use lscratch!!!**
  - **MATLAB batch jobs & swarms should copy the MCR to local scratch!!**
- Don't use /lscratch (or why bother) when...
  - Your job needs to write out a file visible to multiple nodes (this is rare?)
  - Your job is only reading/writing a few files and not very much data.
- **If in doubt, e-mail [staff@hpc.nih.gov](mailto:staff@hpc.nih.gov)**

# Snapshots

- No back-ups of HPC data directories, but there are snapshots.
- A snapshot is a copy of the directory as it existed at a given point in time.
- Current snapshot retention policies:
  - /home – 7 hourly, 6 daily, 8 weekly
  - /data – 2 daily, 2 weekly
  - Weekly snapshots are taken on Sundays for NFS home directories and on Tuesdays for GPFS data directories.

# Accessing snapshots

- Navigate to `.snapshot/` in your `/data` directory
  - Will NOT tab complete; need to type out the full directory name.
  - Users on the VF storage system will need to go to `/vf/users/.snapshot`.
- Will see several subdirectories that contain read-only copies of your directory at a point in time.
- Can copy data from snapshots into your “main” home directory if, for example, you accidentally delete a file!

```
[teacher@helix ~]$ cd /data/teacher
[teacher@helix teacher]$ ls .snapshot
nightly.2020-05-15_0817  nightly.2020-05-18_1005  weekly.2020-05-12_0743  weekly.2020-05-19_0723
[teacher@helix teacher]$ ls .snapshot/nightly.2020-05-15_0817/
convld_predict.py  DL_class1  opt  simplernn_predict.py
convld_train.py    DL_class2  simplernn.h5  simplernn_train.py
```

# Snapshots are NOT back-ups

Snapshots are stored on the same physical hardware as home/data directories. Therefore, **in the event of a hardware or facilities failure, the snapshots will also be lost!**

You **must** back up irretrievable data to a local system (or put it in your /home directory if it will fit).

# A guide to choosing a storage system

- /home
  - Small files that are very important, low I/O intensity
  - Not for files that need to be shared with other users
- /data
  - Bulk data files and scripts
  - Consider using shared data areas for sharing
  - Shared, parallel filesystems – intensive, single node I/O -> /lscratch.
- /scratch
  - **Low I/O intensity** job data that needs to be accessed from multiple nodes.
  - Good for short-term sharing.
- /lscratch
  - **High I/O intensity** job data that only needs to be accessed from a single node.
- Object
  - Primarily read-only data; low to medium intensity, but large capacity.



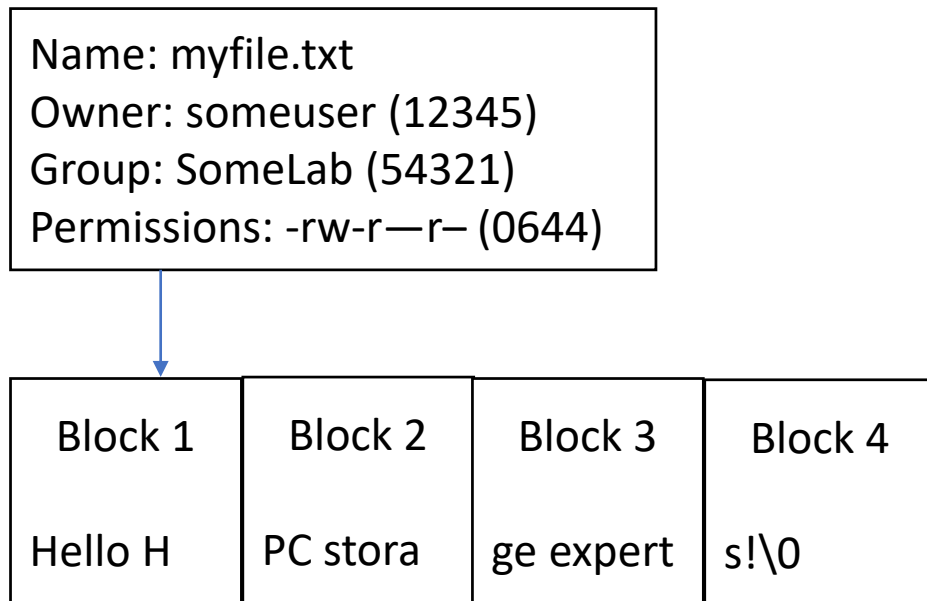


# Presentation overview/outline

- Overview of HPC systems storage
  - Different areas (/home, /data, /scratch, object store)
  - Quotas and quota increases
  - Snapshots
- Understanding input and output (I/O)
  - I/O patterns
  - Introduction to data and metadata
  - A brief look at HPC storage systems
- Putting it all together – using storage effectively
  - Profiling and benchmarking your application's I/O usage
  - Understanding when you're generating too much I/O on the system

# A block about blocks


- A filesystem organizes files and directories into **blocks** of data
  - This is because hard disks can only read and write information in discrete-sized chunks.
- Each file you store that's more than ~ 3 KB takes up at least one full block on the filesystem.
  - The exact size of a block differs from filesystem to filesystem and may be chosen by the administrator who creates the filesystem.




# Sequential vs. random access

- Sequential access: blocks in a file are read one after the other
  - Example: reading in a series of sequences from a file, one after another.
  - Sometimes referred to as “streaming”.
  - E.g. reading a FASTQ file into memory
- Random access: blocks in the file are read in a random order.
  - Common in database applications
  - Much harder for I/O systems to optimize
  - Generally MUCH slower!
  - E.g. Pulling non-adjacent sequences from a BAM file.

## Sequential access



## Random access

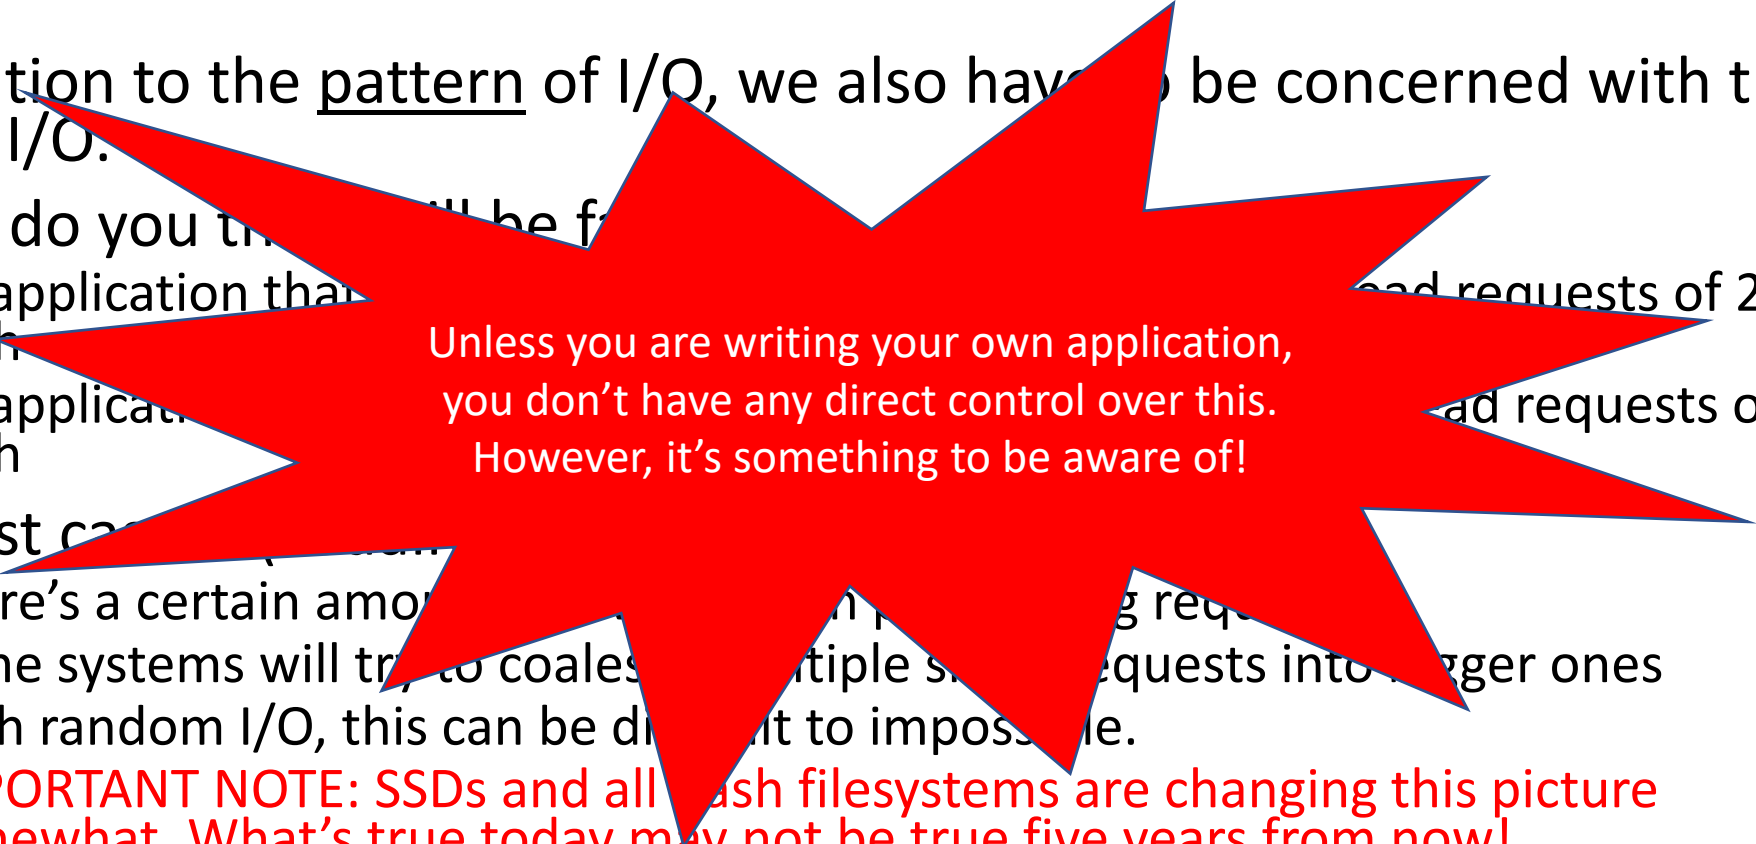


# I/O comes in all sizes

- In addition to the pattern of I/O, we also have to be concerned with the size of I/O.
- Which do you think will be faster?
  - An application that reads 100 MB sequentially by issuing 50 read requests of 2 MB each
  - An application that reads 100 MB sequentially by issuing 5000 read requests of 20 KB each
- The first case is (usually) a lot faster
  - There's a certain amount of overhead in performing requests.
  - Some systems will try to coalesce multiple small requests into bigger ones
  - With random I/O, this can be difficult to impossible.
  - **IMPORTANT NOTE: SSDs and all flash filesystems are changing this picture somewhat. What's true today may not be true five years from now!**

# I/O comes in all sizes


- In addition to the pattern of I/O, we also have to be concerned with the size of I/O.
- Which do you think will be faster?
  - An application that issues 100 read requests of 2 MB each
  - An application that issues 100 read requests of 20 KB each
- The first case is faster
  - There's a certain amount of overhead in processing requests
  - Some systems will try to coalesce multiple small requests into larger ones
  - With random I/O, this can be difficult to impossible.
  - **IMPORTANT NOTE: SSDs and all flash filesystems are changing this picture somewhat. What's true today may not be true five years from now!**



Unless you are writing your own application,  
you don't have any direct control over this.  
However, it's something to be aware of!

# Data vs. metadata


- When we think of a file, we usually think in terms of the **data** it contains.
  - A BAM file contains sequence alignments
  - A molecular dynamics trajectory contains atomic coordinates/velocities
  - An MRI output contains an image of someone's brain.
  - An EM image contains a picture of a cell in the body.
- **Metadata** literally means “data about data”



THINGS GOT TENSE FOR THE TOWNSFOLK  
WHEN THE THIRD META-ANALYST GANG  
RODE INTO TOWN.

# Data vs. metadata

- Examples of metadata
  - When the file was created/accessed
  - The sample ID from which the file was generated
  - The access permissions of the file
  - Where the data is located physically on the underlying disk
  - Can you think of more examples?




THINGS GOT TENSE FOR THE TOWNSFOLK  
WHEN THE THIRD META-ANALYST GANG  
RODE INTO TOWN.



# A conceptual diagram

- Remember our simple schematic...

Name: myfile.txt  
Owner: someuser (12345)  
Group: SomeLab (54321)  
Permissions: -rw-r—r— (0644)



Block 1	Block 2	Block 3	Block 4
Hello H	PC stora	ge expert	s!\0

This is (some of)  
the METADATA

This is the DATA  
in the file.



# Some notes on directories

- Directories are special files that hold pointers (links) to other files.
- The more files there are in a directory, the larger amount of space the directory blocks will take up on disk
- Listing directories, resolving path names, moving files, etc. all require operations on the directory blocks.
  - The file system has to iterate through files in the directory individually
  - The bigger the directory is, the longer these operations will take.
  - This is why the HPC staff recommends having < 5000 files per directory
  - Especially true with directories that will have lots of operations happening simultaneously!

# File and directory parallel access

- Reading and writing the same file from multiple parallel jobs can cause contention.
  - Especially with writing – due to the need for locks to avoid corruption
- Likewise, lots of different processes listing, creating files, etc. in the same directory is a bottleneck.
- Constant creation and deletion of files can create performance issues, particularly when multiple processes are doing it in the same directory.



# Exercise: Good practice or bad practice?

- Which of the following are not good practice? Why?
  - Having 1,000,000 data files in a single directory.
  - Having separate runs of a program write output to separate files.
  - Reading a data file in once at program initiation, and then keeping the data cached in memory.
  - Using the name of a file to encode program results.
  - Having a swarm of 1,000 jobs each use the same temporary directory in /scratch.

# What you can do to avoid bottlenecks!

- Use /lscratch whenever possible!!!
  - Since lscratch is local to the node – avoid all network operations
  - Also, lscratch on newer nodes is provisioned with SSDs!
- Avoid many parallel I/O operations.
  - They tend to oversaturate the disks
- Try to do I/O on large chunks
  - i.e. read and write large amounts of data
  - Less network overhead, and easier for the disk systems to optimize
  - If you're using someone else's code, this is difficult/impossible
- Avoid excessive metadata operations
  - Tend to be filtered to a small amount of disks/controllers.
  - This includes directory operations!

# Exercise – where is the bottleneck?

- For the “bad practices” we identified earlier (marked in red), what bottlenecks are likely to be relevant?
  - Having 1,000,000 data files in a single directory.
  - Having separate runs of a program write output to separate files.
  - Reading a data file in once at program initiation, and then keeping the data cached in memory.
  - Using the name of a file to encode program results.
  - Having a swarm of 1,000 jobs each use the same temporary directory in /scratch.

# Presentation overview/outline

- Overview of HPC systems storage
  - Different areas (/home, /data, /scratch, object store)
  - Quotas and quota increases
  - Snapshots
- Understanding input and output (I/O)
  - I/O patterns
  - Introduction to data and metadata
  - A brief look at HPC storage systems
- Putting it all together – using storage effectively
  - Understanding when you're generating too much I/O on the system

# How do you know if you're abusing the storage systems?

- This can be difficult to know, but there are a few clues
  - Very slow access to working directories involved with the job. (Is etc. take a long time to return)
  - For swarms, job completion speed was acceptable for small numbers of jobs, but gets dramatically slower as the size or number of jobs increases.
    - You did test with small-scale jobs first, right?!
  - The problem may be with another user's jobs, but if you started seeing problems right after you started a bunch of jobs, you are the prime suspect.
- HPC staff will notify you if we notice your jobs having an impact
  - However, please be proactive and don't wait for us to notice the problem
  - If we send you mail, it means you're having a significant negative impact on system performance.

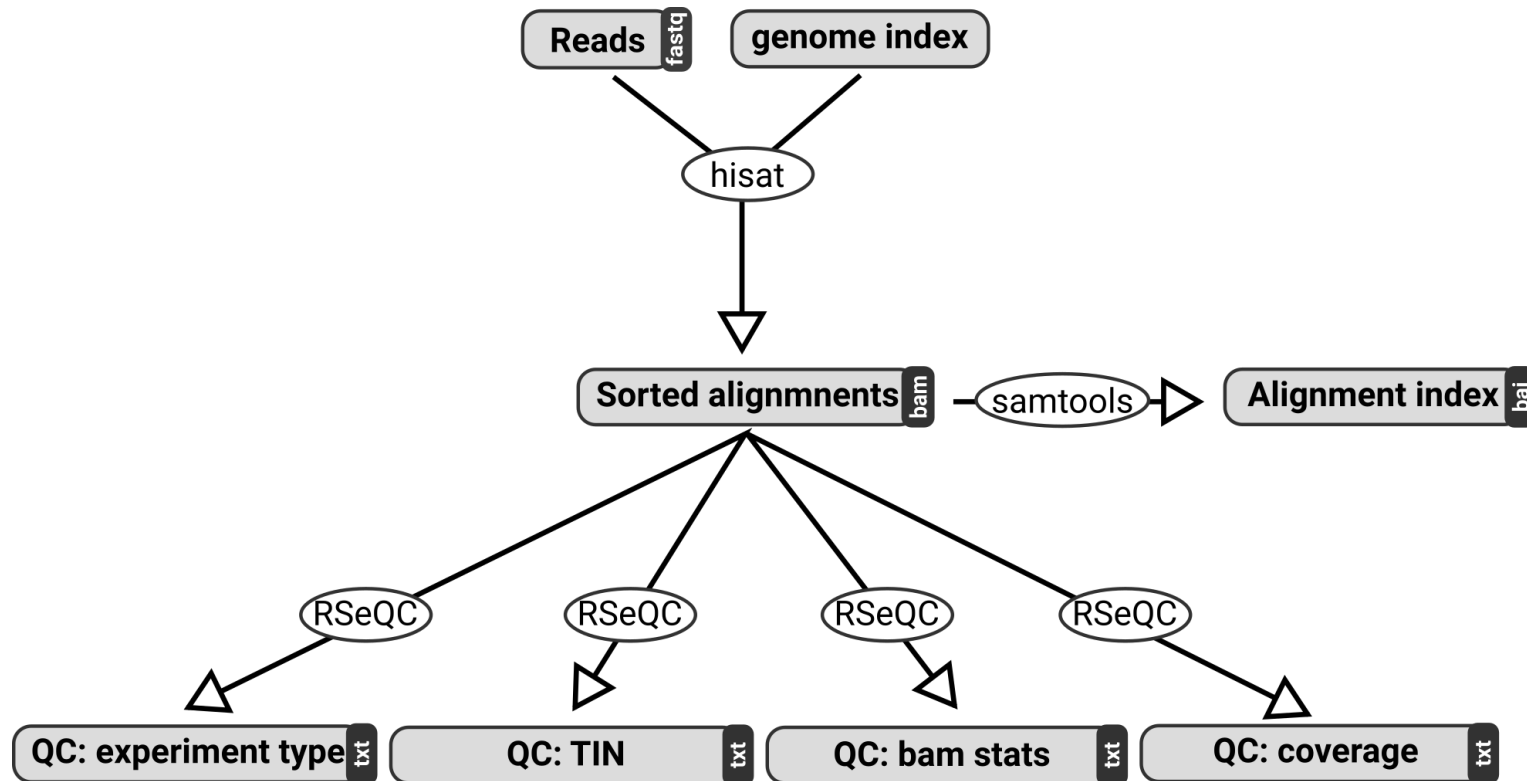
# Refactoring a workload – general principles

- Look for places where lots of parallel processes are doing I/O
  - Think about if only one process could do I/O and communicate with other processes (probably not possible with swarm).
  - Can some or all of that I/O use /lscratch instead of /scratch or /data?
- Think about bottlenecks in the workflow
  - E.g. the whole workload has to wait until one file is updated
  - Does the usage on a shared filesystem cause delays in this process?
- Does the workflow behavior change over time?
  - Do jobs have different I/O patterns in the beginning, middle, or end of their runs.
  - Would staggering this I/O be possible?



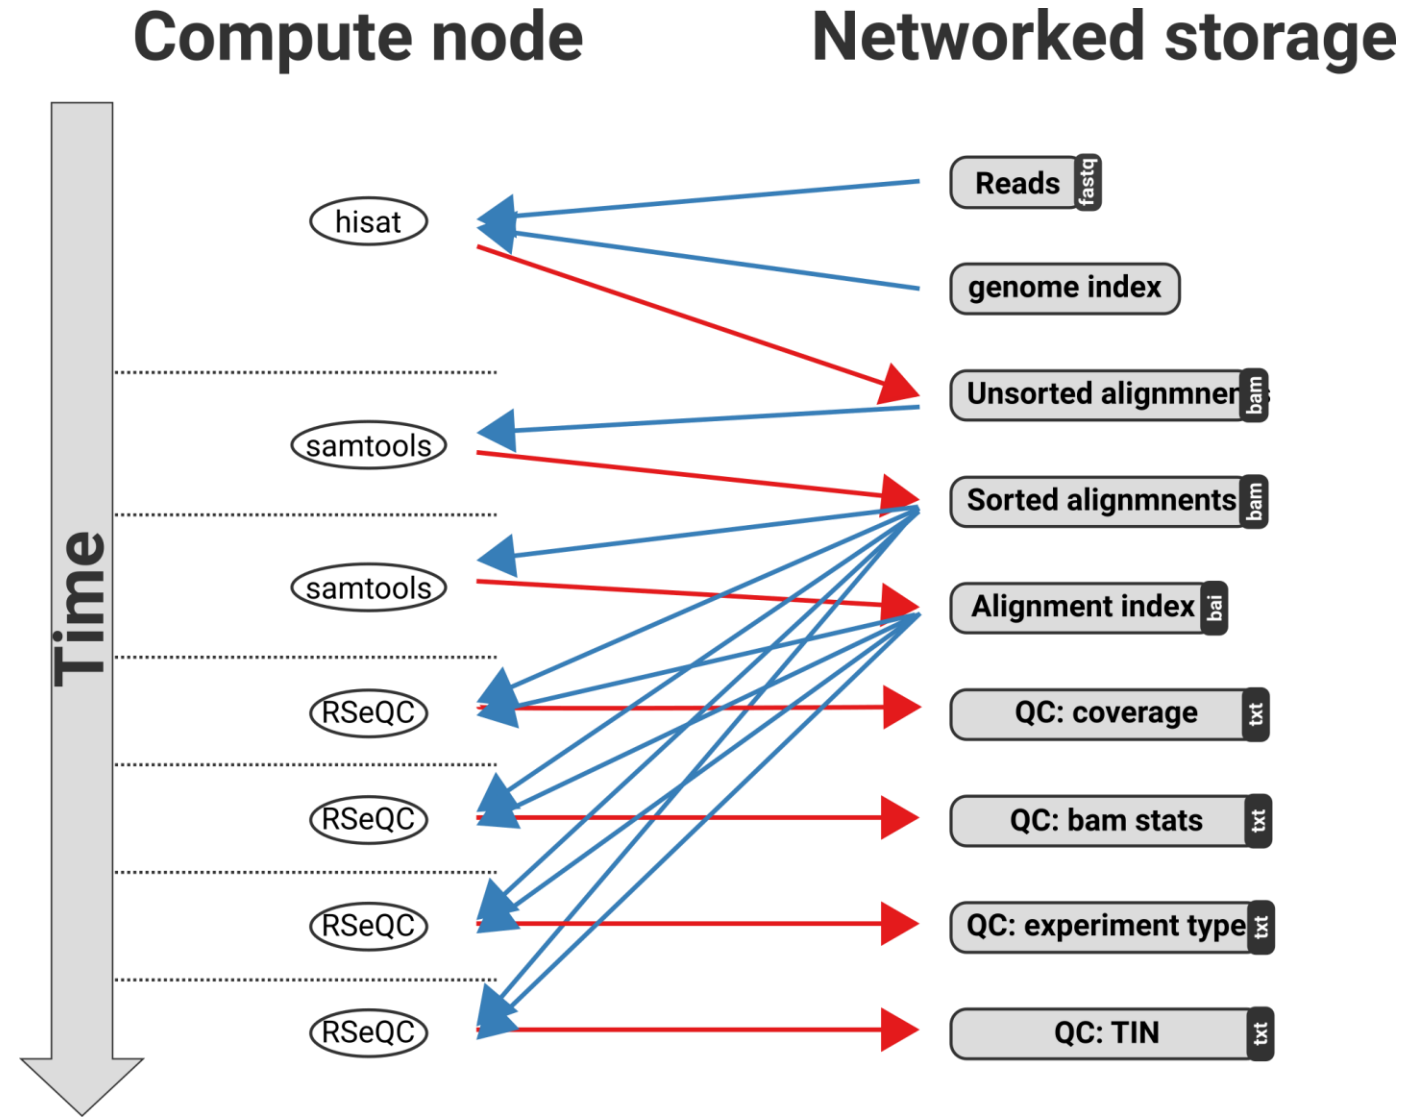
# Workload refactoring: an example

- Start of a real-world genomics pipeline
  - Aligns sequences, creates index
  - RSeQC performs QC steps before continuing the pipeline
- Pipeline ran much more quickly after workload was refactored to use I/O more efficiently!



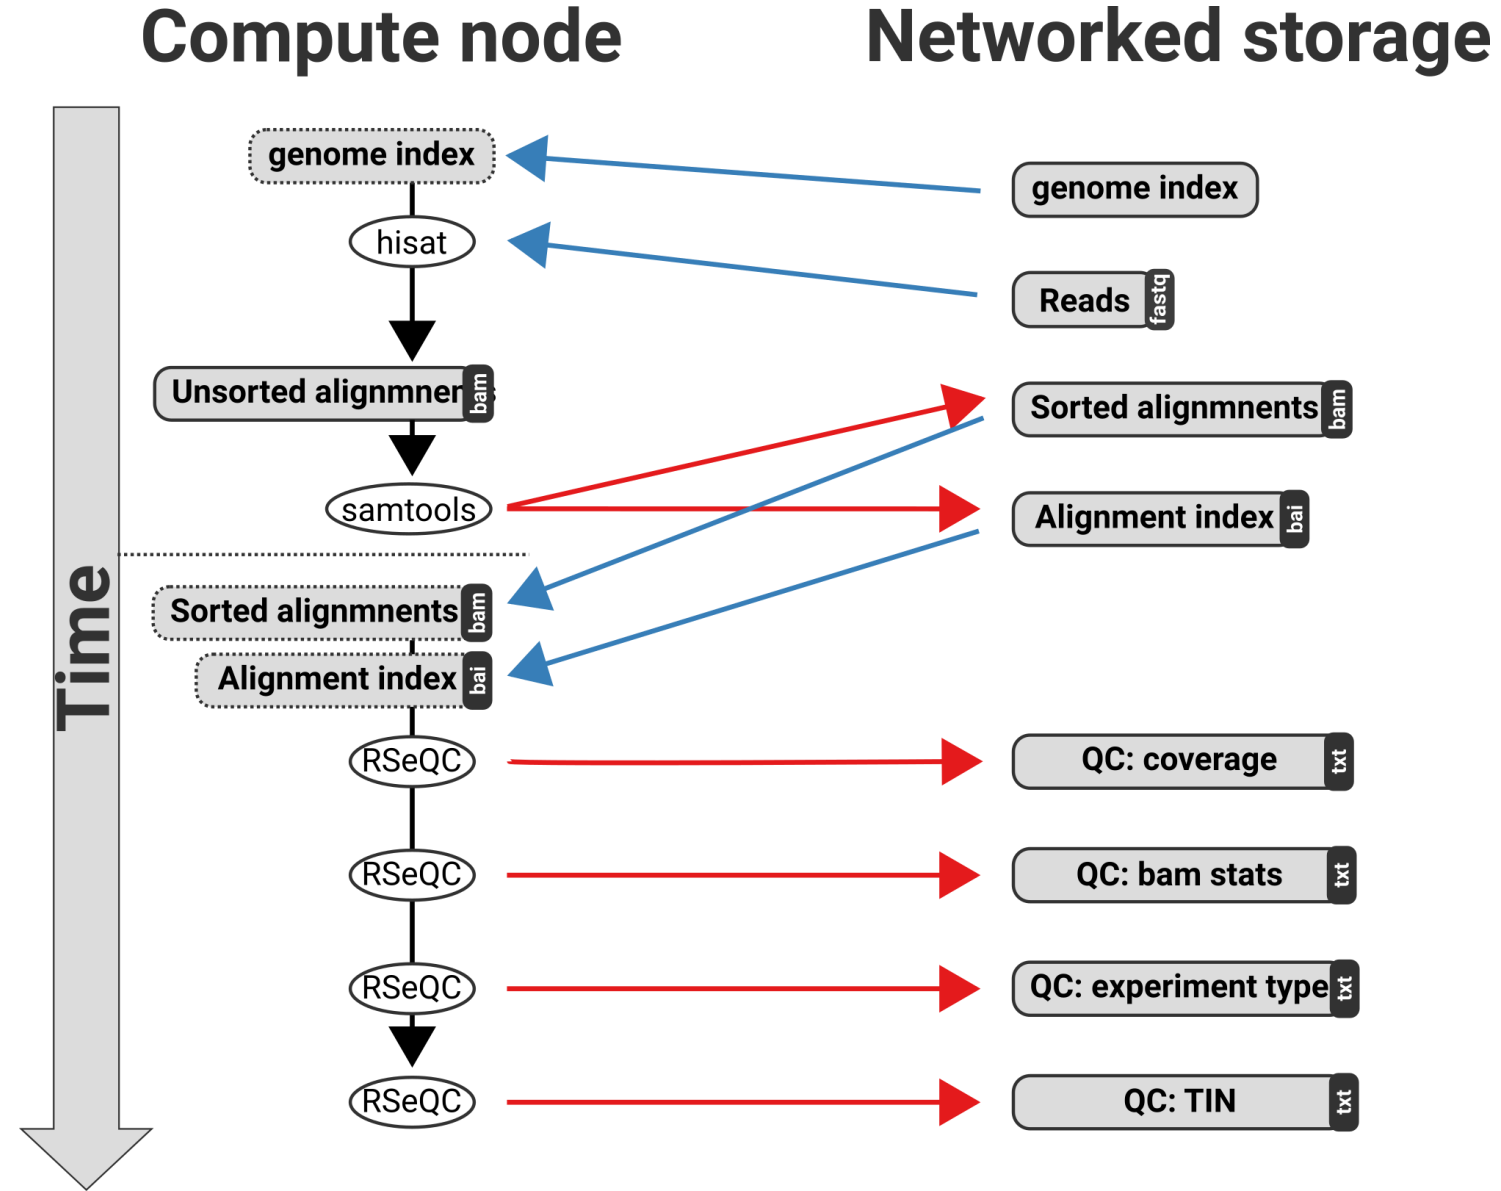
# Original Pipeline

- Lots of read and writes back to network storage.
- Many parallel processes reading the same data from the storage system.
- Depending on the exact analysis done by RSeQC, random I/O is heavy.
  - Remember, lots of RSeQC processes in parallel.
- How would you fix this?




# Refactored pipeline

- Only writes to networked storage when needed.
  - E.g. not after the initial read, only when index is built.
- Instead of each RSeQC reading the data from network storage, use local scratch
- Only write out final result files.



# IOPS comparison

- Bottom line 1: User was able to do more science in less time.
- Bottom line 2: HPC storage admins did not have to troubleshoot performance problems.



# staff@hpc.nih.gov



Steve Bailey



Steven Fellini, Ph.D.



Susan Chacko,  
Ph.D.



Gennady Denisov,  
Ph.D.

**Picture  
unavailable**

Afif Elghraoui



Ali Erfani



Andrew Fant, Ph.D.



Robert Gallimore



David Godlove, Ph.D.



David Hoover, Ph.D.



Patsy Jones

**Picture  
unavailable**

Charles Lehr



Jean Mao, Ph.D.



Tim Miller



Nitish Narula, M. S.



Charlene Osborn



Mark Patkus



Wolfgang Resch,  
Ph.D.



Antonio Ulloa, Ph.D.



Qi Yu, Ph.D.

# Wrap-up; Q&A

- Thank you for coming!
  - We hope you are able to apply the lessons learned to your own particular storage issues.
  - PLEASE reach out to [staff@hpc.nih.gov](mailto:staff@hpc.nih.gov) for assistance; we'd love to work with you **proactively** instead of **reactively**.
- Please provide feedback on this presentation!
  - E-mail Tim [btmiller@hpc.nih.gov](mailto:btmiller@hpc.nih.gov)
  - General questions [staff@hpc.nih.gov](mailto:staff@hpc.nih.gov)

# Basic storage system architecture

- /home and /scratch directories are on a large storage system that uses NFS.
- A few /data directories are on a newer flash based storage system (VF) that also uses NFS.





# Basic storage system architecture

- Other data directories are on systems running IBM's General Parallel File System (GPFS)
- Use “checkquota --gpfs” to determine if any of your data directories are on GPFS.





# NFS and GPFS

- NFS and GPFS have different back-end implementations, but from a user's perspective, they work the same way.
- The systems perform similarly, though file system performance is variable dependent on how many users are accessing a given filesystem at any one time.
  - There is **no** significant performance advantage to using one system vs. the other.
- Main difference from a feature perspective: GPFS and the VF NFS system allow access control lists (ACLs) whereas the NFS implementation used for /home and /scratch does not.
  - ACLs are an advanced way of setting granular file/directory permissions – see <https://hpc.nih.gov/storage/acls.html> for more details
  - We will not discuss ACLs further (unless someone really wants to).

# Figuring out where your data is stored

- You can use the “-a” flag on checkquota to see what filesystems the directories you have access to are on.
- spin1 = NFS, vf = NFS (newer, with ACLs), /gs[4-12] = GPFS
- Can also use “checkquota --gpfs”
- **NEVER** refer to any data directory by its absolute path (i.e. use /data/username **NOT** /gs11/users/username
  - The storage admins move directories for a variety of reasons, so the absolute paths can and do change.